DISCLOSURE DRIFT

Risk disclosures and return forecasting

Research progress report | 10 October 2026

Michael Baffour Awuah | Cornell University

Primary empirical result	Not estimated
Study status	Pipeline implementation and pilot evidence; empirical completion gates remain open

Research question

Do changes in an issuer's annual Risk Factors disclosures improve forecasts of its next 20-session market-adjusted return beyond momentum, volatility, size and historical industry controls?

This study is inspired by Cohen, Malloy and Nguyen's *Lazy Prices* [1]. It narrows the document input to original annual Item 1A sections and fixes a 20-session forecast comparison. It is a new investigation; the source paper's findings are not results for this sample.

Evidence available now

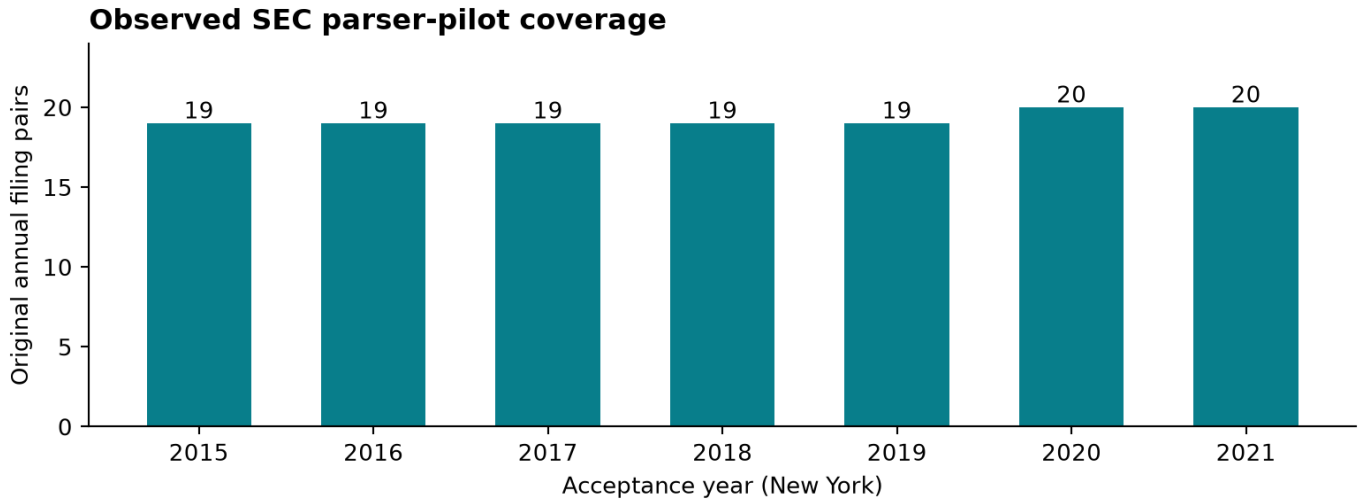
The recovered SEC checkpoint contains **155 original filings, 135 annual pairs and 20 issuers**. The original sources pass the automated audit. Candidate market arithmetic also reproduces, while historical identity, availability, and human reviews remain pending [3-5].

The final 2022-2025 empirical evaluation has not been opened. Figures in this report describe observed SEC text or the frozen method. Any full-pipeline demonstration uses invented data and provides software evidence only.

Public code and audit trail

1. Pilot scope and coverage

The existing sample was selected to exercise document formats, reference exhibits and time conventions. It has no claim to be a representative investment universe. Pair coverage is shown by acceptance year below. The pairing rule preserves issuer identity and consecutive annual filings; predecessor issuers are not silently spliced into the record.



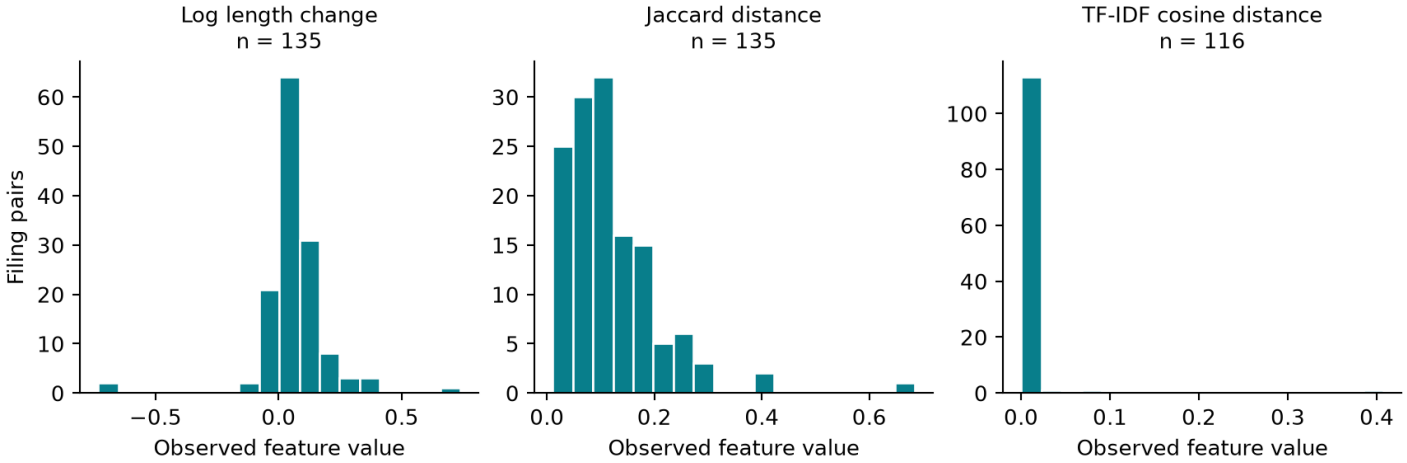
Acceptance year	Annual pairs	Issuers
2015	19	19
2016	19	19
2017	19	19
2018	19	19
2019	19	19
2020	20	20
2021	20	20

The pilot spans 2015-2021; **59** pairs are from 2019-2021. The original plan requested 100-150 pairs in 2019-2021. Coverage over a longer date range does not satisfy that same criterion.

The full study requires a fresh historical population at every year start. The implemented selector uses prior 60-session mean dollar volume, at least 40 valid observations, one most-liquid security per issuer and up to 500 issuers. Later delistings stay in the selected cohort [3].

2. What changed in the text?

Length change measures expansion or contraction in the number of alphabetic tokens. Jaccard distance measures turnover in distinct words. TF-IDF cosine distance weights changes using vocabulary fitted on strictly earlier acceptance years. None is a risk forecast by itself.



Text measure	Observed n	Median	10th / 90th percentile
Current section tokens	135	7601	3448 / 1.268e+04
Log length change	135	0.04595	-0.02304 / 0.1698
Jaccard distance	135	0.09816	0.03452 / 0.2203
TF-IDF cosine distance	116	0.003867	0.00104 / 0.01314

The first 19 pairs have no prior pilot year from which to fit vocabulary, so their TF-IDF values are explicitly unavailable. The remaining 116 distances use earlier years only. This descriptive vocabulary policy does not substitute for the final model's label-maturity filter.

Interpretation limits

Small cosine distances can coexist with meaningful additions to a long section. Token-set distances ignore word order; length does not distinguish a new hazard from added explanation. These summaries justify checking individual passages, not claiming that a model predicts returns.

Example selection is independent of outcomes: the checkpoint selects the smallest SHA-256 of seed and accession within each year, then exposes the source link and a 12-word opening. No largest-return or best-performing example was selected.

3. Timing and audit evidence

For an accepted filing, convert the SEC timestamp to New York time. The reference session is the second XNYS session strictly after its acceptance date. Controls end at the preceding close; the forecast is made before the reference open. Returns begin after the reference close, so the reference day's daily return is excluded [3].

Stage	Evidence and rule
Prior controls	60 exact sessions ending at the preceding close; 20/60 momentum, sample 20-session volatility, log USD capitalization and dated industry
Return label	Product of 1 + daily stock return over the next 20 sessions, minus the same compounded value-weighted market return
Terminal events	Dividends and delisting included once; cash proceeds earn zero after the terminal day while benchmark returns continue
Missing inputs	Unexplained missing sessions, ambiguous security IDs and absent historical snapshots are exclusions or errors; they are never filled with zero

Rebuilt evidence

The source audit records 155 passing filings and 14 registered reference resolutions. It preserves 72 differences between the API-encoded timestamp and original SGML header. The header controls the event date; dissemination time remains an explicit unverified assumption [4].

The matching saved arithmetic receipt records 135 calculated cases and 0 blocked cases, independently reconciled within 1e-12; the maximum absolute difference is 1.43e-15. The fixed 20-trace packet and WFM cash-continuation case remain available. These checks establish arithmetic consistency for candidate mappings, not approval of those mappings [5].

The matched receipts record 0 completed human filing reviews and 0 passing human trace reviews. There are 7 candidate events with a header-CIK difference and 135 links with unresolved open-end markers. These remain M5/M6 completion requirements [5].